

```
In [99]: #Import libraries:
import pandas as pd
import numpy as np
import xgboost as xgb
from xgboost.sklearn import XGBClassifier
from sklearn import model_selection, metrics #Additional sklearn functions
from sklearn.model_selection import GridSearchCV #Performing grid search
import matplotlib.pyplot as plt
%matplotlib inline
```

## Load Data

```
In [100]: train = pd.read_json('../Data/train.json')
```

```
In [101]: train.shape
```

```
Out[101]: (49352, 15)
```

```
In [102]: def calculate_length(df):
            df["num_photos"] = df["photos"].apply(len)
            df["num_features"] = df["features"].apply(len)
            df["num_description_words"] = df["description"].apply(lambda x: len(x.split()))
            return(df)
def calculate_date(df):
    df["created"] = pd.to_datetime(df["created"])
    df["created_year"] = df["created"].dt.year
    df["created_month"] = df["created"].dt.month
    df["created_day"] = df["created"].dt.day
    df["created_hour"] = df["created"].dt.hour
    return(df)
```

Add some additional features

```
In [103]: train = calculate_length(train)
train = calculate_date(train)
train.shape
```

```
Out[103]: (49352, 22)
```

## Label encode certain categorical features

```
In [104]: train.dtypes.unique()
```

```
Out[104]: array([dtype('float64'), dtype('int64'), dtype('O'), dtype('<M8[ns]')], dtype=object)
```

```
In [105]: features = [x for x in train.select_dtypes(include=['int64', 'float64']).columns]
```

```
In [106]: features
```

```
Out[106]: [u'bathrooms',
           u'bedrooms',
           u'latitude',
           u'longitude',
           u'price',
           'num_photos',
           'num_features',
           'num_description_words',
           'created_year',
           'created_month',
           'created_day',
           'created_hour']
```

```
In [107]: train.select_dtypes(include=['object']).columns
```

```
Out[107]: Index([u'building_id', u'description', u'display_address', u'features',
                  u'interest_level', u'manager_id', u'photos', u'street_address'],
                  dtype='object')
```

```
In [108]: from sklearn import preprocessing
```

### Label encode some categorical variables (different from one-hot encoding!!)

```
In [109]: def LabelEncoder(df, columns, features, append=False):
           for cols in columns:
               label = preprocessing.LabelEncoder()
               label.fit(df[cols].values)
               df[cols] = label.transform(df[cols].values)
               if append:
                   features.append(cols)
           return df, features
```

```
In [110]: categorical = ["display_address", "manager_id", "building_id", "street_address"]
           train, features = LabelEncoder(train, categorical, features, True)
```

```
In [111]: train.building_id.head(5)
```

```
Out[111]: 10          2431
          10000       5862
          100004      5806
          100007      1201
          100013         0
          Name: building_id, dtype: int64
```

### Create tf-idf matrix from text features

```
In [112]: train['features'].values
```

```
Out[112]: array([[ ],
                [u'Doorman', u'Elevator', u'Fitness Center', u'Cats Allowed', u'Dogs Allowed'],
                [u'Laundry In Building', u'Dishwasher', u'Hardwood Floors', u'Pets Allowed Case by Case'],
                ...,
                [u'Doorman', u'Elevator', u'Pre-War', u'Dogs Allowed', u'Cats Allowed'],
                [u'Doorman', u'Elevator', u'Pre-War', u'Dogs Allowed', u'Cats Allowed'],
                [u'Hardwood Floors']], dtype=object)
```

```
In [113]: #train['features'] = train["features"].apply(lambda x: " ".join(["_".join(i) for i in x.split(" ")]))
train['features'] = train['features'].apply(lambda x: " ".join(x))
```

```
In [117]: train['features'].values
```

```
Out[117]: array(['', u'Doorman,Elevator,Fitness Center,Cats Allowed,Dogs Allowed',
                u'Laundry In Building,Dishwasher,Hardwood Floors,Pets Allowed Case by Case',
                ..., u'Doorman,Elevator,Pre-War,Dogs Allowed,Cats Allowed',
                u'Doorman,Elevator,Pre-War,Dogs Allowed,Cats Allowed',
                u'Hardwood Floors'], dtype=object)
```

```
In [76]: from sklearn.feature_extraction import text
```

```
In [118]: tfidf = text.CountVectorizer(stop_words='english', max_features=200, ngram_range=(1, 2))
tr_sparse = tfidf.fit_transform(train["features"])
```

```
In [126]: tr_sparse.shape, train[features].shape
```

```
Out[126]: ((49352, 200), (49352, 16))
```

```
In [127]: from scipy import sparse
```

```
In [136]: train_X = sparse.hstack([train[features], tr_sparse]).tocsr()
target_num_map = {'high':0, 'medium':1, 'low':2}
train_y = np.array(train['interest_level'].apply(lambda x: target_num_map[x])).astype(int)
train_X.shape, train_y.shape
```

```
Out[136]: ((49352, 216), (49352,))
```

```
In [155]: np.set_printoptions(threshold=100)
train_X.toarray()[0]
```

```
Out[155]: array([ 1.5,  3., 40.7145, ..., 0., 0., 0.])
```

In [152]: `train[features].head(3)`

Out[152]:

	bathrooms	bedrooms	latitude	longitude	price	num_photos	num_features	num_
10	1.5	3	40.7145	-73.9425	3000	5	0	94
10000	1.0	2	40.7947	-73.9667	5465	11	5	1
100004	1.0	1	40.7388	-74.0018	2850	8	4	93

## Train model

```
In [192]: def xgbfit(model, dtrain, output, useTrainCV=True, cv_folds=5, early_stopping_rounds=10):

    if useTrainCV:
        xgb_param = model.get_xgb_params()
        xgb_param['num_class'] = 3

        xgtrain = xgb.DMatrix(dtrain, label=output)
        cvresult = xgb.cv(xgb_param, xgtrain, num_boost_round=model.get_params()['num_boost_round'],
                           metrics='mlogloss', early_stopping_rounds=early_stopping_rounds)
        model.set_params(n_estimators=cvresult.shape[0])

    #Fit the model algorithm on the data
    model.fit(dtrain, output, eval_metric='mlogloss')

    #Predict training set:
    dtrain_predictions = model.predict(dtrain)

    #Print model report:
    print "\nModel Report"
    #print "R-Square: %.3f" % metrics.r2_score(output, dtrain_predictions)
    print "Log Loss : %.3f" % np.sqrt(metrics.log_loss(output, dtrain_predictions))
    print "Optimal CV Score:"
    print(cvresult.iloc[len(cvresult)-1,:])
    print "Optimal iteration: %d" %(len(cvresult)-1)
    #print "Cross Validation Result: "
    #print(cvresult)

    plt.figure()
    cvresult.loc[:,["test-mlogloss", "train-mlogloss"]].plot()
    return (len(cvresult))
```

```
In [193]: features
```

```
Out[193]: [u'bathrooms',  
           u'bedrooms',  
           u'latitude',  
           u'longitude',  
           u'price',  
           'num_photos',  
           'num_features',  
           'num_description_words',  
           'created_year',  
           'created_month',  
           'created_day',  
           'created_hour',  
           'display_address',  
           'manager_id',  
           'building_id',  
           'street_address']
```

```
In [195]: xgb1 = XGBClassifier(  
           learning_rate=0.1,  
           n_estimators=1000,  
           max_depth=6,  
           min_child_weight=1,  
           #gamma=0,  
           subsample=0.7,  
           colsample_bytree=0.7,  
           objective='multi:softprob',  
           nthread=4,  
           scale_pos_weight=1,  
           seed=189)  
n_estimators = xgbfit(xgb1, train_X, train_y)
```

Model Report

```
In [231]: def runXGB(train_X, train_y, test_X, test_y=None, feature_names=None, seed_val=1234):  
    param = {}  
    param['objective'] = 'multi:softprob'  
    param['eta'] = 0.1  
    param['max_depth'] = 6  
    param['silent'] = True  
    param['num_class'] = 3  
    param['eval_metric'] = "mlogloss"  
    param['min_child_weight'] = 1  
    param['subsample'] = 0.7  
    param['colsample_bytree'] = 0.7  
    param['seed'] = seed_val  
    num_rounds = num_rounds  
  
    plst = list(param.items())  
    xgtrain = xgb.DMatrix(train_X, label=train_y)  
  
    if test_y is not None:  
        xgtest = xgb.DMatrix(test_X, label=test_y)  
        watchlist = [ (xgtrain, 'train'), (xgtest, 'test') ]  
        model = xgb.train(plst, xgtrain, num_rounds, watchlist, early_stopping_rounds=10)  
    else:  
        xgtest = xgb.DMatrix(test_X)  
        model = xgb.train(plst, xgtrain, num_rounds)  
  
    pred_test_y = model.predict(xgtest)  
    return pred_test_y, model
```

```
In [234]: cv_scores = []
kf = model_selection.KFold(n_splits=5, shuffle=True, random_state=9594)
for dev_index, val_index in kf.split(range(train_X.shape[0])):
    # This leads to 5 iterations for 5 splits.
    # dev_index has 80% of the data, val_index has 20% since dev_index takes
    dev_X, val_X = train_X[dev_index:], train_X[val_index:] #training and
    dev_y, val_y = train_y[dev_index], train_y[val_index] #training and val
    preds, model = runXGB(dev_X, dev_y, val_X, val_y, num_rounds=10)
    cv_scores.append(metrics.log_loss(val_y, preds))
print("cv_scores is: \n")
print(cv_scores)
```

```
[0]      train-mlogloss:1.04231  test-mlogloss:1.04335
Multiple eval metrics have been passed: 'test-mlogloss' will be used for
early stopping.
```

Will train until test-mlogloss hasn't improved in 30 rounds.

```
[1]      train-mlogloss:0.988808  test-mlogloss:0.991342
[2]      train-mlogloss:0.9447    test-mlogloss:0.947949
[3]      train-mlogloss:0.905032  test-mlogloss:0.909496
[4]      train-mlogloss:0.873738  test-mlogloss:0.879315
[5]      train-mlogloss:0.8452    test-mlogloss:0.851965
[6]      train-mlogloss:0.821483  test-mlogloss:0.828979
[7]      train-mlogloss:0.798081  test-mlogloss:0.806763
[8]      train-mlogloss:0.779337  test-mlogloss:0.789054
[9]      train-mlogloss:0.763191  test-mlogloss:0.773825
[0]      train-mlogloss:1.04218   test-mlogloss:1.04306
```

Multiple eval metrics have been passed: 'test-mlogloss' will be used for early stopping.

Will train until test-mlogloss hasn't improved in 30 rounds.

```
[1]      train-mlogloss:0.98889   test-mlogloss:0.990589
[2]      train-mlogloss:0.944869  test-mlogloss:0.947613
[3]      train-mlogloss:0.905808  test-mlogloss:0.909032
[4]      train-mlogloss:0.874403  test-mlogloss:0.878286
[5]      train-mlogloss:0.846063  test-mlogloss:0.85052
[6]      train-mlogloss:0.82223   test-mlogloss:0.82755
[7]      train-mlogloss:0.798562  test-mlogloss:0.804737
[8]      train-mlogloss:0.780015  test-mlogloss:0.78687
[9]      train-mlogloss:0.764084  test-mlogloss:0.771618
[0]      train-mlogloss:1.03727   test-mlogloss:1.03705
```

Multiple eval metrics have been passed: 'test-mlogloss' will be used for early stopping.

Will train until test-mlogloss hasn't improved in 30 rounds.

```
[1]      train-mlogloss:0.984484  test-mlogloss:0.984005
[2]      train-mlogloss:0.944027  test-mlogloss:0.943723
[3]      train-mlogloss:0.905473  test-mlogloss:0.905309
[4]      train-mlogloss:0.872217  test-mlogloss:0.872441
[5]      train-mlogloss:0.84519   test-mlogloss:0.845638
[6]      train-mlogloss:0.821019  test-mlogloss:0.821717
[7]      train-mlogloss:0.797477  test-mlogloss:0.798822
[8]      train-mlogloss:0.777528  test-mlogloss:0.779398
[9]      train-mlogloss:0.758671  test-mlogloss:0.761119
[0]      train-mlogloss:1.03672   test-mlogloss:1.03845
```

Multiple eval metrics have been passed: 'test-mlogloss' will be used for early stopping.

Will train until test-mlogloss hasn't improved in 30 rounds.

```
[1]    train-mlogloss:0.984133 test-mlogloss:0.986989
[2]    train-mlogloss:0.943447 test-mlogloss:0.94748
[3]    train-mlogloss:0.904404 test-mlogloss:0.909973
[4]    train-mlogloss:0.871004 test-mlogloss:0.877333
[5]    train-mlogloss:0.843771 test-mlogloss:0.850999
[6]    train-mlogloss:0.819389 test-mlogloss:0.827512
[7]    train-mlogloss:0.795706 test-mlogloss:0.804846
[8]    train-mlogloss:0.77611 test-mlogloss:0.785978
[9]    train-mlogloss:0.757212 test-mlogloss:0.767794
[0]    train-mlogloss:1.03689 test-mlogloss:1.03828
```

Multiple eval metrics have been passed: 'test-mlogloss' will be used for early stopping.

Will train until test-mlogloss hasn't improved in 30 rounds.

```
[1]    train-mlogloss:0.983946 test-mlogloss:0.986643
[2]    train-mlogloss:0.943248 test-mlogloss:0.946835
[3]    train-mlogloss:0.904048 test-mlogloss:0.909118
[4]    train-mlogloss:0.870739 test-mlogloss:0.877103
[5]    train-mlogloss:0.843342 test-mlogloss:0.850627
[6]    train-mlogloss:0.818933 test-mlogloss:0.82725
[7]    train-mlogloss:0.795661 test-mlogloss:0.805011
[8]    train-mlogloss:0.776003 test-mlogloss:0.78638
[9]    train-mlogloss:0.756996 test-mlogloss:0.768503
```

cv\_scores is:

```
[0.77382506840606557, 0.77161818340597343, 0.76111855259299155, 0.7677942
3060448582, 0.7685032073548016]
```